



太原理工大学
TAIYUAN UNIVERSITY OF TECHNOLOGY



太原理工大学
大数据学院
COLLEGE OF DATA SCIENCE
TAIYUAN UNIVERSITY OF TECHNOLOGY

第二章 深度前馈网络

深度前馈网络

主要内容

01 CNN反向传播

02 深度前馈网络

03 TensorFlow

PART CNN反向传播 ONE

CNN反向传播


全链接：前向传播与反向回馈

Square Euclidean Distance (regression)

$$J = \frac{1}{2} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

$$J(W, b; x, y) = \frac{1}{2} \|h_{W,b}(x) - y\|^2$$

$$= \left[\frac{1}{m} \sum_{i=1}^m \left(\frac{1}{2} \|h_{W,b}(x^{(i)}) - y^{(i)}\|^2 \right) \right] + \frac{\lambda}{2} \sum_{l=1}^{m-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} (W_{ji}^{(l)})^2$$



更新迭代:

$$W_{ij}^{(l)} = W_{ij}^{(l)} - \alpha \frac{\partial}{\partial W_{ij}^{(l)}} J(W, b)$$

$$b_i^{(l)} = b_i^{(l)} - \alpha \frac{\partial}{\partial b_i^{(l)}} J(W, b)$$

$$\frac{\partial}{\partial W_{ij}^{(l)}} J(W, b) = \left[\frac{1}{m} \sum_{i=1}^m \frac{\partial}{\partial W_{ij}^{(l)}} J(W, b; x^{(i)}, y^{(i)}) \right] + \lambda W_{ij}^{(l)}$$

$$\frac{\partial}{\partial b_i^{(l)}} J(W, b) = \frac{1}{m} \sum_{i=1}^m \frac{\partial}{\partial b_i^{(l)}} J(W, b; x^{(i)}, y^{(i)})$$

CNN反向传播

全链接：反向推导

假设神经网络(NN)总共有 L 层

当第 $L-1$ 层时，权重求导

$$\frac{\partial J}{\partial W_{ij}^{L-1}} = \frac{\partial J}{\partial z_i^L} \frac{\partial z_i^L}{\partial W_{ij}^{L-1}} = \delta_i^L a_j^{L-1}$$

sigmoid函数 $f'(z) = f(z)(1 - f(z))$

tanh函数 $f'(z) = 1 - (f(z))^2$

$$\delta_i^L = \frac{\partial J}{\partial z_i^L} = \frac{\partial}{\partial z_i^L} \sum_{i=1}^{s_L} \frac{1}{2} \|y_i - f(z_i^L)\|^2 = -(y_i - f(z_i^L)) f'(z_i^L)$$

当第 $L-2$ 层时，权重求导

$$\frac{\partial J}{\partial W_{ij}^{L-2}} = \frac{\partial J}{\partial z_i^{L-1}} \frac{\partial z_i^{L-1}}{\partial W_{ij}^{L-2}} = \delta_i^{L-1} a_j^{L-2}$$

$$\begin{aligned} \delta_i^{L-1} &= \frac{\partial J}{\partial z_i^{L-1}} = \frac{\partial}{\partial z_i^{L-1}} \sum_{b=1}^{s_L} \frac{1}{2} \|y_b - f(z_b^L)\|^2 = \sum_{b=1}^{s_L} -(y_b - f(z_b^L)) f'(z_b^L) \frac{\partial z_b^L}{\partial z_i^{L-1}} \\ &= \sum_{b=1}^{s_L} \delta_b^L \cdot w_{bi}^{L-1} f'(z_i^{L-1}) \\ &= \left(\sum_{b=1}^{s_L} \delta_b^L w_{bi}^{L-1} \right) f'(z_i^{L-1}) \end{aligned}$$










5

CNN反向传播

http://galaxy.agh.edu.pl/~vlsi/AL/backp_t_en/backprop.html

全链接：反向推导





正向

6


CNN反向传播

全链接：反向推导



$$\frac{\partial J}{\partial w_{ij}^{L-1}} = -(y_i - f(z_i^L))f'(z_i^L)a_j^{L-1}$$

$$\frac{\partial J}{\partial w_{ij}^{L-2}} = \left(\sum_{k=1}^{s_L} \delta_k^L w_{ki}^{L-1} \right) f'(z_i^{L-1}) a_j^{L-2}$$








误差反向传导

7


CNN反向传播

全链接：反向推导



$$\frac{\partial J}{\partial w_{ij}^{L-1}} = -(y_i - f(z_i^L)) f'(z_i^L) a_j^{L-1}$$

$$\frac{\partial J}{\partial w_{ij}^{L-2}} = \left(\sum_{k=1}^{s_L} \delta_k^L w_{ki}^{L-1} \right) f'(z_i^{L-1}) a_j^{L-2}$$




梯度更新

8

CNN神经网络

CNN: 训练 (前向传播)



```

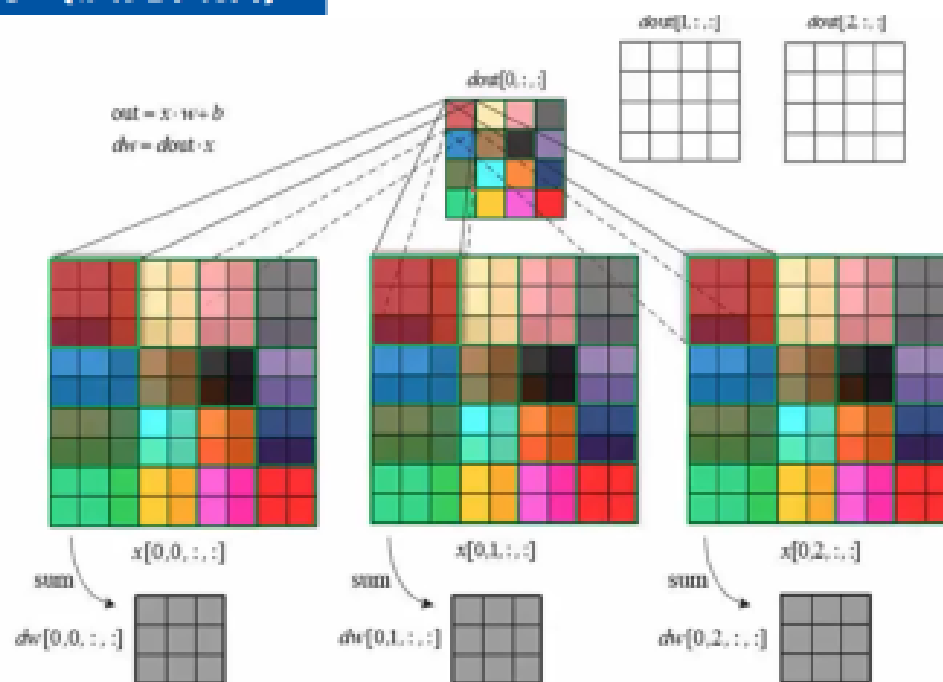
out[0,0,0] = np.sum(x[0, :, : 3, : 3] * w[0]) + b[0]
out[0,0,1] = np.sum(x[0, :, 2 : 5, : 3] * w[0]) + b[0]
...
out[1,0,0] = np.sum(x[0, :, : 3, : 3] * w[1]) + b[1]
out[1,0,1] = np.sum(x[0, :, 2 : 5, : 3] * w[1]) + b[1]
...
out[2,0,0] = np.sum(x[0, :, : 3, : 3] * w[2]) + b[2]
out[2,0,1] = np.sum(x[0, :, 2 : 5, : 3] * w[2]) + b[2]
...

```



CNN神经网络


CNN: 训练 (反向传播)



$$\begin{aligned}
 dw[0,0,:,:] &= x[0,0,3,3] * dout[0,0,0] + x[0,0,3,2:5] * dout[0,0,1] + \dots + x[0,0,6:9,6:9] * dout[0,3,3] \\
 dw[0,1,:,:] &= x[0,1,3,3] * dout[0,0,0] + x[0,1,3,2:5] * dout[0,0,1] + \dots + x[0,1,6:9,6:9] * dout[0,3,3] \\
 dw[0,2,:,:] &= x[0,2,3,3] * dout[0,0,0] + x[0,2,3,2:5] * dout[0,0,1] + \dots + x[0,2,6:9,6:9] * dout[0,3,3]
 \end{aligned}$$

CNN反向传播

推导



forward: [1 3; 2 2] -> [2]

backward: [2] -> [0.5 0.5; 0.5 0.5]

forward: [1 3; 2 2] -> 3

backward: [3] -> [0 3; 0 0]

PART 深度前馈网络 TWO

深度前馈网络

通用近似定理

定理 4.1 – 通用近似定理 (Universal Approximation Theorem) [Cybenko, 1989, Hornik et al., 1989]: 令 $\varphi(\cdot)$ 是一个非常数、有界、单调递增的连续函数, \mathcal{I}_d 是一个 d 维的单位超立方体 $[0, 1]^d$, $C(\mathcal{I}_d)$ 是定义在 \mathcal{I}_d 上的连续函数集合。对于任何一个函数 $f \in C(\mathcal{I}_d)$, 存在一个整数 m , 和一组实数 $v_i, b_i \in \mathbb{R}$ 以及实数向量 $\mathbf{w}_i \in \mathbb{R}^d, i = 1, \dots, m$, 以至于我们可以定义函数

$$F(\mathbf{x}) = \sum_{i=1}^m v_i \varphi(\mathbf{w}_i^T \mathbf{x} + b_i), \quad (4.33)$$

作为函数 f 的近似实现, 即

$$|F(\mathbf{x}) - f(\mathbf{x})| < \epsilon, \forall \mathbf{x} \in \mathcal{I}_d. \quad (4.34)$$

其中 $\epsilon > 0$ 是一个很小的正数。


根据通用近似定理, 对于具有线性输出层和至少一个使用“挤压”性质的激活函数的隐藏层组成的前馈神经网络, 只要其隐藏层神经元的数量足够, 它可以以任意的精度来近似任何从一个定义在实数空间中的有界闭集函数。

深度前馈网络

通用近似定理

- Given a **shallow** network structure with one hidden layer with ReLU activation and linear output

A piece-wise linear functions



- Given a L-Lipschitz function f^*
 - How many neurons are needed to approximate f^* ?

深度前馈网络

通用近似定理

- Given a L -Lipschitz function f^*
 - How many neurons are needed to approximate f^* ?

L -Lipschitz Function (smooth)

$$\|f(x_1) - f(x_2)\| \leq L \|x_1 - x_2\|$$

Output
change

Input
change

$L=1$ for "1-Lipschitz"



深度前馈网络

通用近似定理

$$\begin{aligned} \max_{0 \leq x \leq 1} |f(x) - f^*(x)| &\leq \varepsilon \\ \downarrow \\ \sqrt{\int_0^1 |f(x) - f^*(x)|^2 dx} &\leq \varepsilon \end{aligned}$$

- Given a L-Lipschitz function f^*
 - How many neurons are needed to approximate f^* ?


$f \in N(K)$ \longrightarrow The function space defined by the network with K neurons.

Given a small number $\varepsilon > 0$

What is the number of K such that

$$\text{Exist } f \in N(K), \max_{0 \leq x \leq 1} |f(x) - f^*(x)| \leq \varepsilon$$

The difference between $f(x)$ and $f^*(x)$ is smaller than ε .



深度前馈网络

通用近似定理


Universality

- L-Lipschitz function f^*

All the functions in $N(K)$ are piecewise linear.

Approximate f^* by a piecewise linear function f

How to make the errors $\leq \varepsilon$




深度前馈网络

通用近似定理

Universality

- L-Lipschitz function f^*

How to make a 1 hidden layer relu network have the output like green curve?



深度前馈网络

通用近似定理

▶ 1. 实现

- ▶ 使用Numpy实现前馈神经网络

▶ 2. 函数拟合

- ▶ 理论和实验证明，一个两层的ReLU网络可以模拟任何函数

- ▶ https://github.com/nndl/exercise/tree/master/for_chapter_4_%20simple%20neural%20network

Exponential Representation Advantage of Depth




Figure 6.5

PART **TensorFlow** THREE

<http://playground.tensorflow.org/>

Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.



TensorFlow PlayGround



- ❑ 选择Sigmoid函数作为激活函数，明显能感觉到训练的时间很长，ReLU函数能大大加快收敛速度
- ❑ 当把隐含层数加深后，会发现Sigmoid函数作为激活函数，训练过程loss降不下来
- ❑ 隐含层的数量不是越多越好，层数和特征的个数太多，会造成优化的难度和出现过拟合的现象
- ❑ 只需要输入最基本的特征 x_1 , x_2 ，只要给予足够多层的神经网络和神经元，神经网络会自己组合出最有用的特征

TensorFlow PlayGround

Epoch
001,892

Learning rate
0.00

Activation
ReLU

Regularization
None

Regularization rate
0

Problem type
Classification



24

TensorFlow

安装

1. Windows



CPU版:

环境: python 3.5, 3.6(64位)

本地pip安装:

```
pip3 install --upgrade tensorflow
```

Anaconda安装:

(1) 创建一个名为tensorflow的conda环境

```
conda create -n tensorflow pip python = 3.5
```

(2) 激活conda

```
activate tensorflow环境
```

(3) 在conda环境中安装TensorFlow

```
pip install --ignore-installed --upgrade tensorflow
```

TensorFlow

安装

1. Windows



GPU版：

环境

- (1) python3.5及以上（64位）
- (2) GPU卡: 计算力不小于3.0的NVIDIA显卡
- (3) CUDA工具包9.0
- (4) cuDNN v7.0

2. Ubuntu（Ubuntu 16.04或更高版本）

环境与window要求一致

https://tensorflow.google.cn/install/install_linux



3. macOS（macOS X 10.11（El Capitan）或更高版本）

https://tensorflow.google.cn/install/install_mac




26

TensorFlow

Basic concepts

- Express a numeric computation as a **graph**.
- Graph nodes are **operations** which have any number of inputs and outputs
- Graph edges are **tensors** which flow between nodes

$$h_i = \text{ReLU}(Wx + b)$$



27


TensorFlow

Basic concepts

$$h_i = \text{ReLU}(Wx + b)$$

Variables are 0-ary stateful nodes which output their current value. (State is retained across multiple executions of a graph.)

(parameters, gradient stores, eligibility traces, ...)




TensorFlow

Basic concepts

$$h_i = \text{ReLU}(Wx + b)$$

Placeholders are 0-ary nodes whose value is fed in at execution time.

(inputs, variable learning rates, ...)



TensorFlow

Basic concepts


$$h_i = \text{ReLU}(Wx + b)$$

Mathematical operations:

MatMul: Multiply two matrix values.

Add: Add elementwise (with broadcasting).

ReLU: Activate with elementwise rectified linear function.



TensorFlow

Basic concepts

1. Create model weights,
including initialization


$a.W \sim \text{Uniform}(-1, 1); b = 0$

2. Create input placeholder x

$a.m * 784$ input matrix

3. Create computation
graph

$$h_i = \text{ReLU}(Wx + b)$$



```
import tensorflow as tf
```

```
1  b = tf.Variable(tf.zeros((100,)))  
   W = tf.Variable(tf.random_uniform((784, 100),  
                                     -1, 1))  
2  x = tf.placeholder(tf.float32, (None, 784))  
3  h_i = tf.nn.relu(tf.matmul(x, W) + b)
```

Just Run It!

```
sess = tf.Session()  
sess.run(tf.initialize_all_variables())  
sess.run(h_i, {x: np.random.random(64, 784)})
```

31

TensorFlow

Basic concepts

1. Build a graph

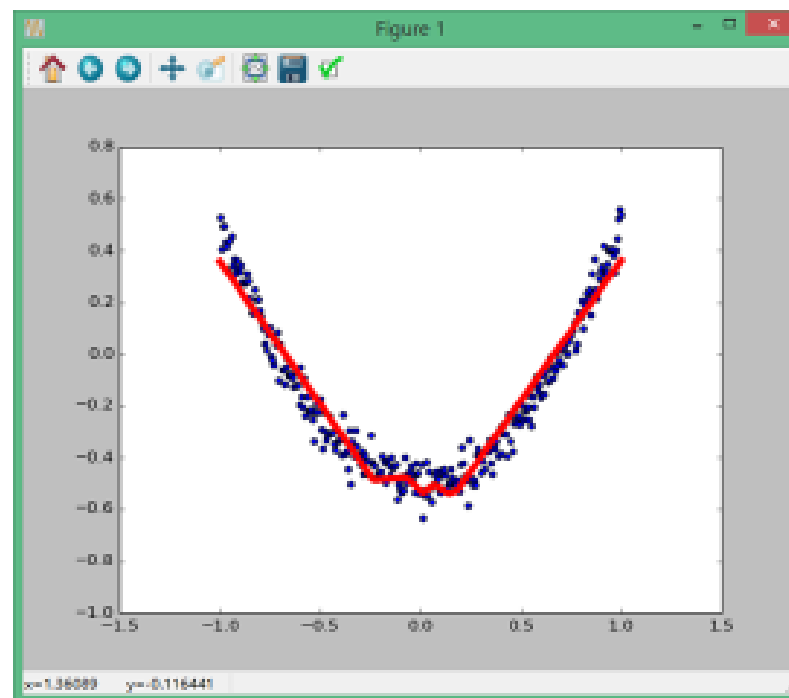
- a. Graph contains parameter specifications, model architecture, optimization process, ...
- b. Somewhere between 5 and 5000 lines

2. Initialize a session

3. Fetch and feed data with `Session.run`

- a. Compilation, optimization, etc. happens at this step
 - you probably won't notice


TensorFlow TOY



33

深度前馈网络

单层



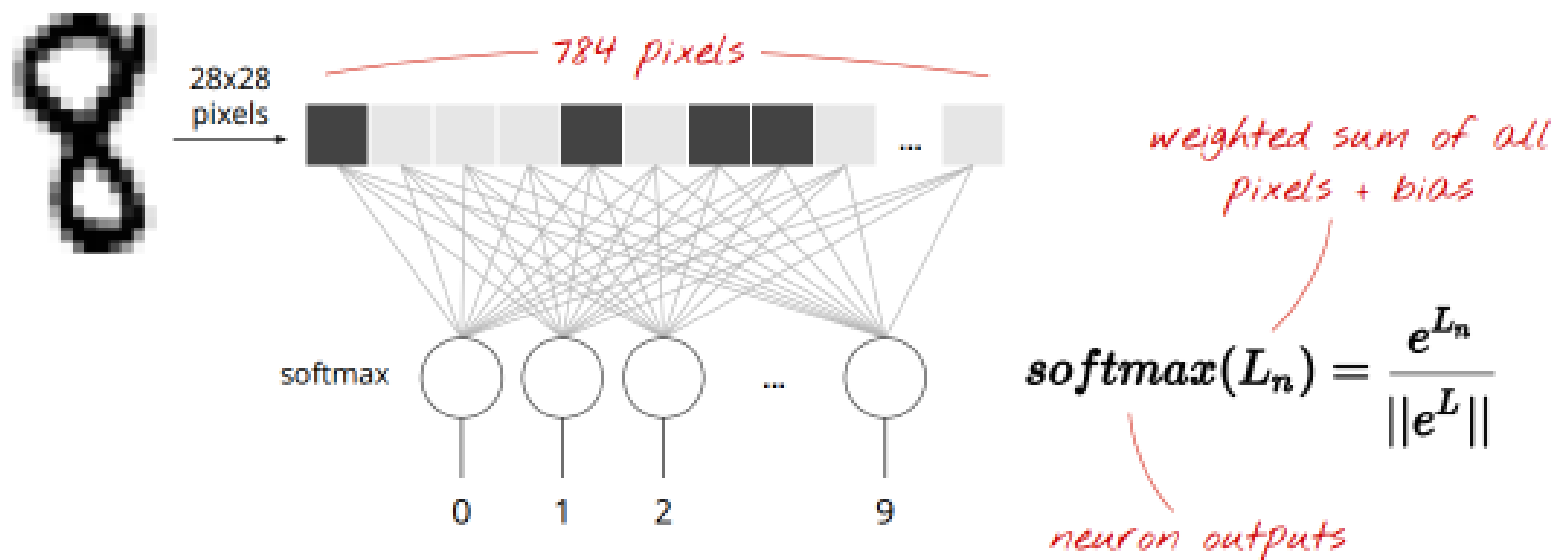
MNIST = Mixed National Institute of Standards and Technology - Download the dataset at <http://yann.lecun.com/exdb/mnist/>

34

深度前馈网络

单层

$$h_{\theta}(x) = \begin{bmatrix} P(y=1|x;\theta) \\ P(y=2|x;\theta) \\ \vdots \\ P(y=K|x;\theta) \end{bmatrix} = \frac{1}{\sum_{j=1}^K \exp(\theta^{(j)\top} x)} \begin{bmatrix} \exp(\theta^{(1)\top} x) \\ \exp(\theta^{(2)\top} x) \\ \vdots \\ \exp(\theta^{(K)\top} x) \end{bmatrix}$$



<http://deeplearning.stanford.edu/tutorial/supervised/SoftmaxRegression/>

35

深度前馈网络

单层

$$h_{\theta}(x) = \begin{bmatrix} P(y=1|x;\theta) \\ P(y=2|x;\theta) \\ \vdots \\ P(y=K|x;\theta) \end{bmatrix} = \frac{1}{\sum_{j=1}^K \exp(\theta^{(j)\top} x)} \begin{bmatrix} \exp(\theta^{(1)\top} x) \\ \exp(\theta^{(2)\top} x) \\ \vdots \\ \exp(\theta^{(K)\top} x) \end{bmatrix}$$

当类别数 $k = 2$ 时
softmax 回归退化为 logistic 回归

When $K=2$, the softmax regression hypothesis outputs

$$h_{\theta}(x) = \frac{1}{\exp(\theta^{(1)\top} x) + \exp(\theta^{(2)\top} x)} \begin{bmatrix} \exp(\theta^{(1)\top} x) \\ \exp(\theta^{(2)\top} x) \end{bmatrix}$$



$$\begin{aligned} h(x) &= \frac{1}{\exp((\theta^{(1)} - \theta^{(2)})^{\top} x) + \exp(\theta^{(2)\top} x)} \begin{bmatrix} \exp((\theta^{(1)} - \theta^{(2)})^{\top} x) \exp(\theta^{(2)\top} x) \\ \exp(\theta^{(2)\top} x) \end{bmatrix} \\ &= \begin{bmatrix} \frac{1}{1 + \exp((\theta^{(1)} - \theta^{(2)})^{\top} x)} \\ \frac{\exp((\theta^{(1)} - \theta^{(2)})^{\top} x)}{1 + \exp((\theta^{(1)} - \theta^{(2)})^{\top} x)} \end{bmatrix} \\ &= \begin{bmatrix} \frac{1}{1 + \exp((\theta^{(1)} - \theta^{(2)})^{\top} x)} \\ 1 - \frac{1}{1 + \exp((\theta^{(1)} - \theta^{(2)})^{\top} x)} \end{bmatrix} \end{aligned}$$

36

深度前馈网络

单层

tensor shapes: $X[100, 748]$ $W[748, 10]$ $b[10]$

$Y = \text{tf.nn.softmax}(\text{tf.matmul}(X, W) + b)$

matrix multiply

broadcast on all lines

37

深度前馈网络

单层

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
| 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |

actual probabilities, "one-hot" encoded

Cross entropy: $-\sum Y_i' \cdot \log(Y_i)$

computed probabilities

| | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.1 | 0.2 | 0.1 | 0.3 | 0.2 | 0.1 | 0.9 | 0.2 | 0.1 | 0.1 |
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |

this is a "6"

深度前馈网络

单层

```
import tensorflow as tf
```

```
X = tf.placeholder(tf.float32, [None, 28, 28, 1])
```

```
W = tf.Variable(tf.zeros([784, 10]))
```

```
b = tf.Variable(tf.zeros([10]))
```

```
init = tf.initialize_all_variables()
```

this will become the batch size, 100

28 x 28 grayscale images

Training = computing variables W and b

深度前馈网络

单层

```
# model
Y = tf.nn.softmax(tf.matmul(tf.reshape(X, [-1, 784]), W) + b)
# placeholder for correct answers
Y_ = tf.placeholder(tf.float32, [None, 10])

# loss function
cross_entropy = -tf.reduce_sum(Y_ * tf.log(Y))

# % of correct answers found in batch
is_correct = tf.equal(tf.argmax(Y, 1), tf.argmax(Y_, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

flattening images

"one-hot" encoded

"one-hot" decoding

40

深度前馈网络

单层

```
optimizer = tf.train.GradientDescentOptimizer(0.003)
train_step = optimizer.minimize(cross_entropy)
```

learning rate

loss function

41

深度前馈网络

单层

```
sess = tf.Session()
sess.run(init)

for i in range(1000):
    # Load batch of images and correct answers
    batch_X, batch_Y = mnist.train.next_batch(100)
    train_data={X: batch_X, Y_: batch_Y}

    # train
    sess.run(train_step, feed_dict=train_data)

    # success ?
    a,c = sess.run([accuracy, cross_entropy], feed_dict=train_data)

    # success on test data ?
    test_data={X: mnist.test.images, Y_: mnist.test.labels}
    a,c = sess.run([accuracy, cross_entropy, It], feed=test_data)
```

running a Tensorflow computation, feeding placeholders

Tip: do this every 100 iterations

42

深度前馈网络

单层

```
import tensorflow as tf
```

initialisation

```
X = tf.placeholder(tf.float32, [None, 28, 28, 1])  
W = tf.Variable(tf.zeros([784, 10]))  
b = tf.Variable(tf.zeros([10]))  
init = tf.initialize_all_variables()
```

model

```
Y=tf.nn.softmax(tf.matmul(tf.reshape(X,[-1, 784]), W) + b)
```

placeholder for correct answers

```
Y_ = tf.placeholder(tf.float32, [None, 10])
```

Loss function

```
cross_entropy = -tf.reduce_sum(Y_ * tf.log(Y))
```

success metrics

% of correct answers found in batch

```
is_correct = tf.equal(tf.argmax(Y,1), tf.argmax(Y_,1))  
accuracy = tf.reduce_mean(tf.cast(is_correct,tf.float32))
```

```
optimizer = tf.train.GradientDescentOptimizer(0.003)  
train_step = optimizer.minimize(cross_entropy)
```

training step

```
sess = tf.Session()  
sess.run(init)
```

```
for i in range(10000):
```

Load batch of images and correct answers

```
batch_X, batch_Y = mnist.train.next_batch(100)
```

```
train_data={X: batch_X, Y_: batch_Y}
```

train

```
sess.run(train_step, feed_dict=train_data)
```

Run

success ? add code to print it

```
a,c = sess.run([accuracy, cross_entropy], feed=train_data)
```

success on test data ?

```
test_data={X:mnist.test.images, Y_:mnist.test.labels}
```

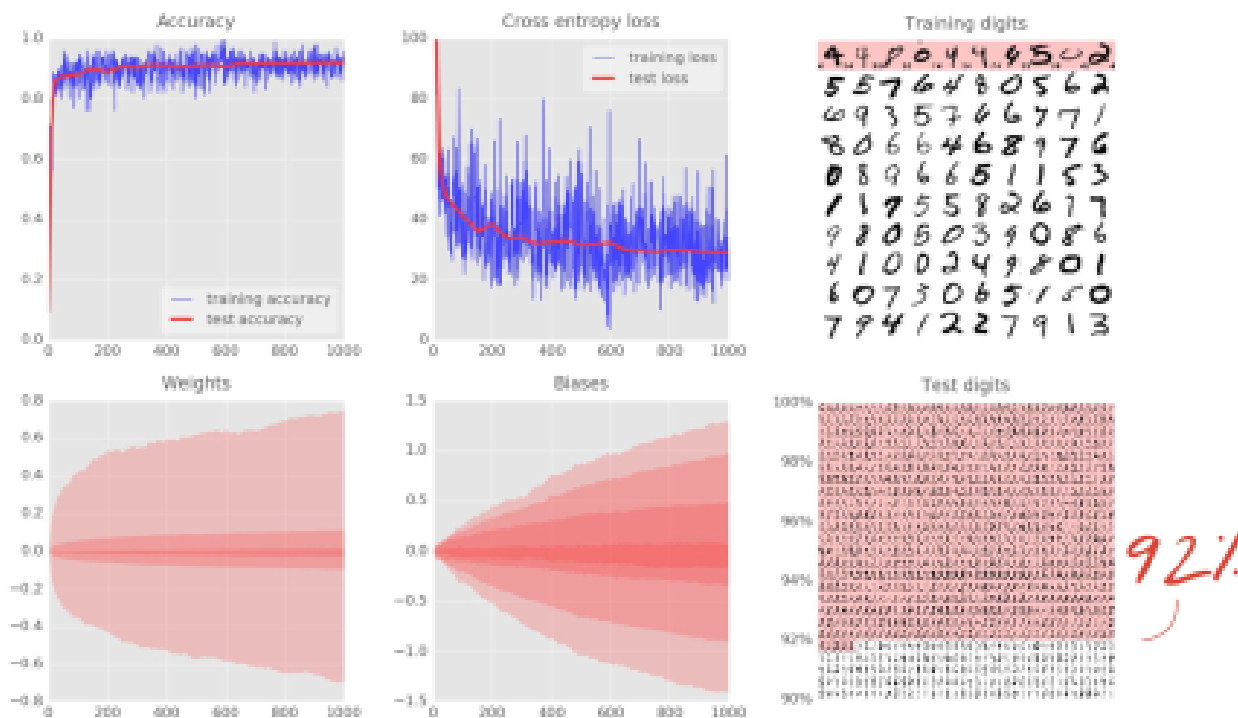
```
a,c = sess.run([accuracy, cross_entropy], feed=test_data)
```

43

DEMO


深度前馈网络

单层



深度前馈网络



多层



深度前馈网络

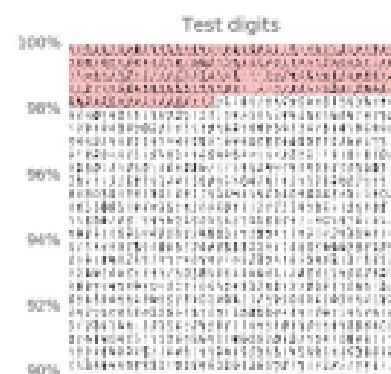


多层

Too many neurons



Training digits

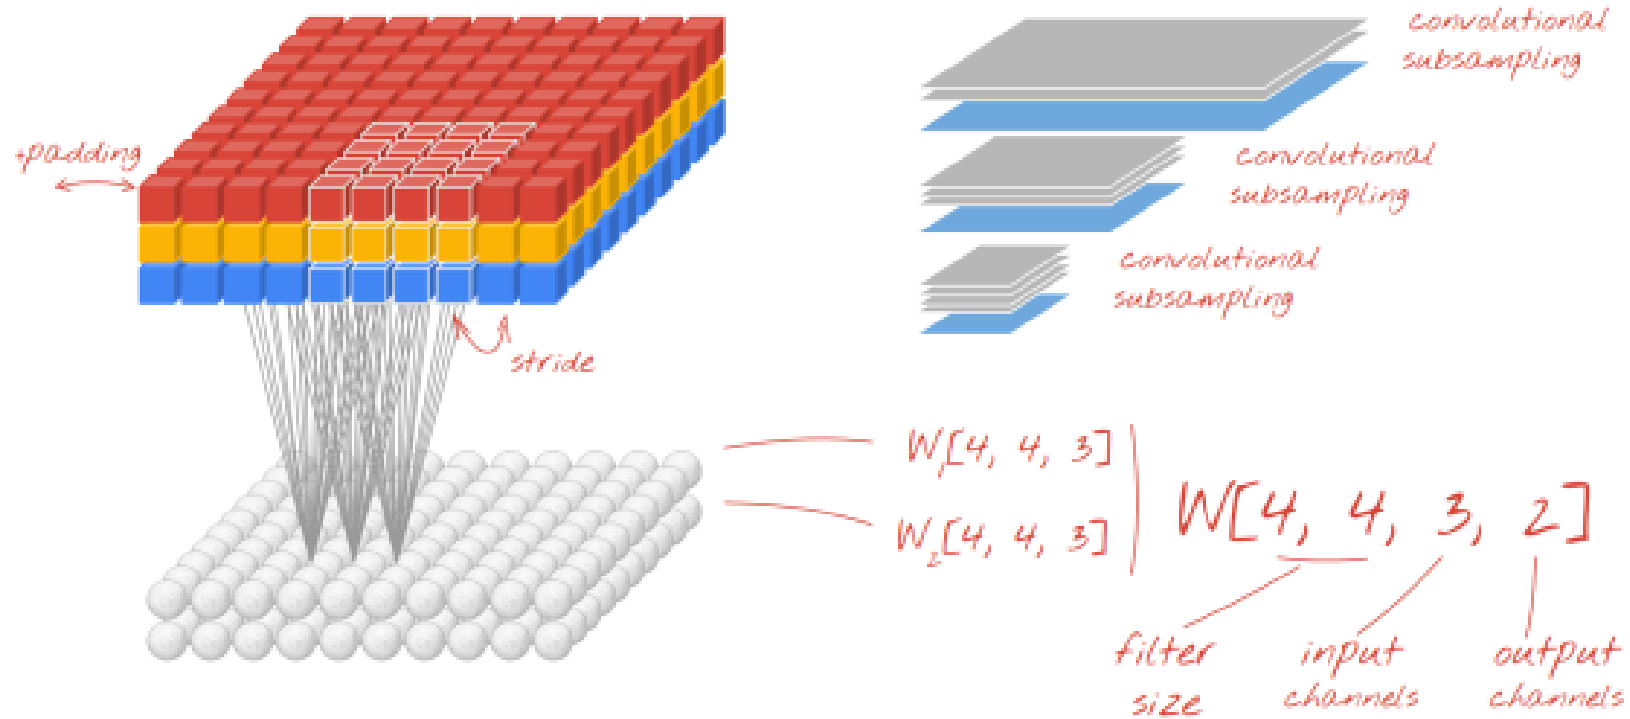
4071034180
5008360123
6003185349
1925802010
8483984258
0614538551
3276015854
8662091045
9964825151
0581127990




98%

深度前馈网络


卷积





Less parameters!

Even less parameters!



深度前馈网络

卷积



DEMO

50

深度前馈网络

卷积层

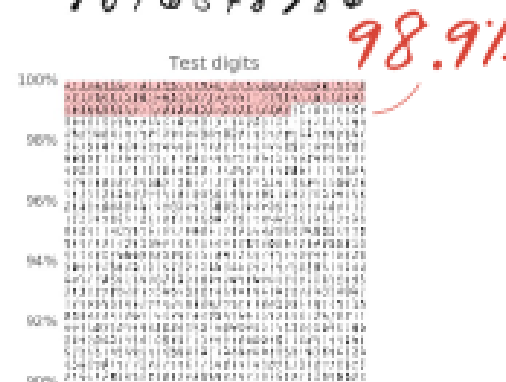


Still Too many neurons



Training digits



```

7 4 9 6 2 9 1 6 1 1
1 0 6 6 6 6 1 3 3 3
7 0 7 2 6 1 9 5 4 6
4 4 2 8 9 5 2 7 5 2
7 0 9 0 4 5 0 6 5 4
3 1 3 2 6 2 5 2 7 7
1 2 1 7 0 6 8 0 3 9
5 9 9 9 7 2 7 7 6 5
9 9 2 8 7 1 3 7 4 4
1 0 1 6 5 7 8 7 8 6
    
```





深度前馈网络

Bigger卷积+ dropout



Training digits

9 4 1 9 2 1 6 9 4 9
0 1 6 1 5 3 8 1 9 3
9 6 8 9 9 0 1 6 2 8
0 1 2 6 3 3 1 1 5 8
4 3 2 4 5 2 0 0 0 9
3 3 4 7 5 0 5 6 1 4
7 9 7 3 6 4 1 9 5 2
9 0 1 8 0 4 5 6 8 2
4 3 0 1 0 0 5 7 2 5
1 2 4 4 2 8 5 6 7 0



Test digits

9 4 1 9 2 1 6 9 4 9
0 1 6 1 5 3 8 1 9 3
9 6 8 9 9 0 1 6 2 8
0 1 2 6 3 3 1 1 5 8
4 3 2 4 5 2 0 0 0 9
3 3 4 7 5 0 5 6 1 4
7 9 7 3 6 4 1 9 5 2
9 0 1 8 0 4 5 6 8 2
4 3 0 1 0 0 5 7 2 5
1 2 4 4 2 8 5 6 7 0

99.3%

Excellent

THANK YOU
Q&A

